

## 2. Scaling Model

Perform capacity planning for a chat application that supports 40,000 concurrent users.

Estimate the network, storage, and compute requirements.

Use simple, back-of-the-napkin math to justify the scaling model.

**Deliverables:**

**Diagrams & Documentation: Present findings in a markdown file, spreadsheet,**

**Perform capacity planning for a chat application that supports 40,000 concurrent users.**

Capacity planning for a chat application with 40,000 concurrent users requires a comprehensive assessment of various factors. Here's a breakdown of the process:

### 1. Define Application Requirements:

- **Message Rate:**
  - Estimate the average number of messages per user per minute or second.
  - Consider peak usage periods.
- **Message Size:**
  - Determine the average message size (text, images, files).
- **Features:**
  - Identify features that impact resource usage (e.g., group chats, file sharing, video calls, presence indicators).
- **Data Storage:**
  - Estimate the amount of data to be stored (messages, user profiles, files).
  - Consider data retention policies.
- **Latency Requirements:**
  - Define acceptable latency for message delivery and other operations.
- **Availability Requirements:**
  - Determine the desired uptime (e.g., 99.9%, 99.99%).

### 2. Estimate Resource Requirements:

- **Server Resources:**
  - **CPU:**
    - Message processing, real-time communication, and application logic.
    - Consider using load testing to determine CPU usage under peak load.
  - **Memory (RAM):**
    - User sessions, message buffers, and application data.
    - Estimate memory usage per user and multiply by the number of concurrent users.
  - **Network Bandwidth:**
    - Message traffic, file transfers, and real-time communication.

- Calculate bandwidth based on message rate and size.
- **Storage:**
  - Message storage, user profiles, and file storage.
  - Estimate storage requirements based on data volume and retention policies.
- **Database Resources:**
  - **Database Servers:**
    - Message storage, user profiles, and other application data.
    - Consider using a scalable database solution (e.g., NoSQL, cloud-based database).
  - **Database Storage:**
    - Estimate database storage requirements based on data volume and retention policies.
    - Database IOPS.
- **Real-time Communication Resources:**
  - **WebSockets/SignalR Servers:**
    - Handle real-time communication between users.
    - Estimate the number of connections and message traffic.
  - **Load Balancers:**
    - Distribute traffic across multiple servers.
    - Ensure high availability and scalability.

### 3. Capacity Planning Calculations (Example):

- **Assumptions:**
  - Average message rate: 5 messages per user per minute.
  - Average message size: 1 KB.
  - 40,000 concurrent users.
- **Calculations:**
  - Total messages per minute:  $40,000 \text{ users} * 5 \text{ messages/user} = 200,000$  messages.
  - Total data per minute:  $200,000 \text{ messages} * 1 \text{ KB/message} = 200 \text{ MB}$ .
  - Total data per second:  $200\text{MB}/60 = \sim 3.3\text{MB/s}$
  - Bandwidth requirements depend on the connection type and overhead.
- **Server Calculations:**
  - These calculations are highly dependant on the software used, and the architecture of the application. Load testing is required.
  - Example: if each user requires 1MB of ram, then 40GB of ram is required.
  - CPU usage testing is required.

### 4. Technology Choices:

- **Cloud Platforms:** AWS, Azure, GCP (provide scalability and elasticity).
- **Databases:**
  - NoSQL databases (e.g., Cassandra, MongoDB) for high-volume message storage.
  - Relational databases (e.g., PostgreSQL, MySQL) for user profiles and other structured data.
  - Cloud-based databases (e.g., Azure Cosmos DB, AWS DynamoDB) for scalability and availability.

- **Real-time Communication:**
  - WebSockets.
  - SignalR.
  - Socket.IO.
- **Load Balancing:**
  - Cloud-based load balancers (e.g., Azure Load Balancer, AWS ELB).
  - Nginx, HAProxy.
- **Caching:**
  - Redis, Memcached.

## 5. Load Testing:

- **Simulate Peak Load:**
  - Use load testing tools (e.g., JMeter, Locust, Gatling) to simulate 40,000 concurrent users and their message traffic.
- **Measure Performance:**
  - Monitor CPU usage, memory usage, network bandwidth, and latency during load testing.
- **Identify Bottlenecks:**
  - Identify performance bottlenecks and optimize application and infrastructure.
- **Scale Gradually:**
  - Increase the number of concurrent users gradually to determine the application's breaking point.

## 6. Monitoring and Scaling:

- **Real-time Monitoring:**
  - Implement real-time monitoring of application and infrastructure metrics.
- **Auto-Scaling:**
  - Configure auto-scaling to automatically adjust resources based on demand.
- **Alerting:**
  - Set up alerts for critical metrics (e.g., high CPU usage, high latency).

## Important Considerations:

- **Architecture:** The application's architecture plays a significant role in its scalability. Consider using a microservices architecture for better scalability.
- **Optimization:** Optimize database queries, message processing, and network communication for performance.
- **Redundancy:** Implement redundancy to ensure high availability and fault tolerance.
- **Security:** Implement security measures to protect user data and prevent unauthorized access.
- **Geographic Distribution:** If your users are geographically dispersed, consider using a content delivery network (CDN) and deploying your application in multiple regions.

By following these steps, you can create a robust capacity plan for your chat application and ensure that it can handle 40,000 concurrent users.

# Estimate the network, storage, and compute requirements.

Let's break down the estimation of network, storage, and compute requirements for a chat application supporting 40,000 concurrent users, building upon the previous capacity planning outline.

## Assumptions (For Example Calculations):

- Average message rate: 5 messages per user per minute.
- Average message size: 1 KB.
- 40,000 concurrent users.
- Data retention: 1 year.
- Average user profile size: 5 KB.
- Average file upload per user per day: 1 file, 1 MB in size.
- Peak traffic is 3 times the average.

## 1. Network Requirements:

- **Average Message Traffic:**
  - Total messages per minute:  $40,000 \text{ users} * 5 \text{ messages/user} = 200,000$  messages.
  - Total data per minute:  $200,000 \text{ messages} * 1 \text{ KB/message} = 200 \text{ MB}$ .
  - Total data per second:  $200 \text{ MB} / 60 \text{ seconds} = \sim 3.33 \text{ MB/s}$ .
  - Peak Traffic:  $3.33 \text{ MB/s} * 3 = \sim 10 \text{ MB/s}$ .
- **File Transfer Traffic:**
  - Total file uploads per day:  $40,000 \text{ users} * 1 \text{ MB/user} = 40,000 \text{ MB}$  (40 GB).
  - This will be spread throughout the day, so it will not add the full amount to the peak traffic.
- **Real-time Communication Overhead:**
  - WebSockets/SignalR connections introduce overhead.
  - Estimate 20% overhead for control messages and keep-alive signals.
- **Total Network Bandwidth:**
  - Consider peak traffic, real-time overhead, and file transfer bursts.
  - A minimum of 1 Gbps (125 MB/s) bandwidth is recommended. This allows for headroom.
  - For a cloud environment, ensure that the virtual machines or container instances have sufficient network throughput.

## 2. Storage Requirements:

- **Message Storage:**
  - Total messages per year:  $200,000 \text{ messages/minute} * 60 \text{ minutes/hour} * 24 \text{ hours/day} * 365 \text{ days/year} = \sim 105 \text{ billion messages}$ .
  - Total message storage:  $105 \text{ billion messages} * 1 \text{ KB/message} = \sim 105 \text{ TB}$ .
- **User Profile Storage:**
  - Total user profile storage:  $40,000 \text{ users} * 5 \text{ KB/user} = 200 \text{ MB}$ . This is very small.
- **File Storage:**
  - Total file storage per day: 40 GB.
  - Total file storage per year:  $40 \text{ GB/day} * 365 \text{ days/year} = \sim 14.6 \text{ TB}$ .
- **Database Storage:**
  - Include indexes, logs, and other database overhead.

- Add a 20-30% buffer to message and user profile storage.
- **Total Storage:**
  - Approximately 120-130 TB (including buffer and overhead).
  - Consider object storage for files, and a scalable database for messages.

### 3. Compute Requirements:

- **Application Servers:**
  - CPU:
    - Message processing, real-time communication, and application logic.
    - Load testing is essential to determine precise CPU requirements.
    - Start with a cluster of medium-sized VMs or containers (e.g., 8 vCPUs each).
    - Use auto-scaling to adjust based on CPU utilization.
  - Memory:
    - User sessions, message buffers, and application data.
    - Estimate memory usage per user and multiply by the number of concurrent users.
    - Start with 16-32 GB RAM per server.
- **Database Servers:**
  - CPU:
    - Database operations (queries, indexing).
    - Use a database cluster with sufficient CPU power.
  - Memory:
    - Database caching and operations.
    - Allocate sufficient RAM for the database.
  - Database IOPS:
    - The database will be very write heavy. Plan for high IOPS.
- **Real-time Communication Servers (WebSockets/SignalR):**
  - CPU:
    - Connection management and message routing.
    - Use a cluster of servers with load balancing.
  - Memory:
    - Connection state and message buffers.
    - Allocate sufficient RAM for connection management.
- **Load Balancers:**
  - Distribute traffic across application and real-time communication servers.
  - Use cloud-based load balancers for scalability and availability.

### Important Notes:

- **Load Testing:** Load testing is crucial to validate these estimates and identify performance bottlenecks.
- **Cloud Scalability:** Cloud platforms provide elasticity, allowing you to scale resources as needed.
- **Monitoring:** Implement comprehensive monitoring to track resource utilization and performance.
- **Optimization:** Optimize application code, database queries, and network communication for performance.
- **Redundancy:** Implement redundancy to ensure high availability and fault tolerance.
- **Caching:** Use caching to reduce database load and improve performance.

- **Microservices:** Consider using a microservices architecture for better scalability and maintainability.
- **Geographic Distribution:** If your users are geographically dispersed, consider using a CDN and deploying your application in multiple regions.

## Use simple, back-of-the-napkin math to justify the scaling model

Alright, let's do some quick, back-of-the-napkin calculations to justify our scaling model for the chat app with 40,000 concurrent users.

### 1. Message Traffic (Network):

- **Users:** 40,000
- **Messages/User/Minute:** 5
- **Total Messages/Minute:**  $40,000 * 5 = 200,000$
- **Message Size:** 1 KB
- **Data/Minute:**  $200,000 \text{ KB} = 200 \text{ MB}$
- **Data/Second:**  $200 \text{ MB} / 60 \text{ seconds} \approx 3.3 \text{ MB/s}$
- **Peak (3x Average):**  $3.3 \text{ MB/s} * 3 \approx 10 \text{ MB/s}$

#### Justification:

- Even at peak, 10 MB/s isn't a huge amount. But, we need to account for overhead and bursts.
- A 1 Gbps (125 MB/s) connection gives us ample headroom, ensuring smooth performance even during spikes.

### 2. Storage:

- **Messages/Year:**  $200,000 \text{ messages/minute} * 60 * 24 * 365 \approx 105 \text{ billion messages}$
- **Message Storage:**  $105 \text{ billion KB} \approx 105 \text{ TB}$
- **File Uploads/Day:**  $40,000 \text{ users} * 1 \text{ MB} \approx 40 \text{ GB}$
- **File Storage/Year:**  $40 \text{ GB/day} * 365 \text{ days} \approx 14.6 \text{ TB}$

#### Justification:

- Message storage will dominate. 105 TB is a lot, so we need a scalable storage solution (like cloud object storage).
- File storage adds another 15 TB, so we're looking at roughly 120-130 TB total.

### 3. Compute (Simplified):

- **Users:** 40,000
- **"User Load Unit":** Let's say, for simplicity, each user requires a small, consistent amount of compute, a "User Load Unit" (ULU). This ULU represents a combination of CPU and RAM.
- **Server Capacity:** A modest server can handle, say, 1,000 ULUs.
- **Total ULUs:**  $40,000 \text{ users} * 1 \text{ ULU/user} = 40,000 \text{ ULUs}$
- **Servers Needed:**  $40,000 \text{ ULUs} / 1,000 \text{ ULUs/server} = 40 \text{ servers (roughly)}$

**Justification:**

- This is a very rough estimate. The actual ULU value depends on the application's efficiency.
- This shows we'll need a cluster of servers, not just one.
- We need to add more servers for redundancy and peak loads.
- We also need to have separate servers for the database, and real time communication.

**Key points:**

- **Network:** We need sufficient bandwidth to handle peak message traffic and bursts.
- **Storage:** We need a scalable storage solution to accommodate the large volume of messages and files.
- **Compute:** We need a cluster of servers to handle the load, with auto-scaling to adjust based on demand.
- **Load Testing is Key:** These are just rough estimates. Load testing is essential to validate and refine our scaling model.
- **Cloud Benefits:** Cloud platforms offer the elasticity and scalability we need to handle this scale.